

Automation with Ansible

Version 0.1.1

April 2014

Automation with Ansible

Version 0.1.1

April 2014

Adham Helal
Engin Yöyen

Copyright © 2014 Adham Helal, Engin Yöyen.

Permission is granted to copy, distribute, and/or modify this document under the terms of the Creative Commons Attribution-NonCommercial 3.0 Unported License, which is available at

<http://creativecommons.org/licenses/by-nc/3.0/>.

The original form of this book is \LaTeX source code. Compiling this \LaTeX source has the effect of generating a device-independent representation of a textbook, which can be converted to other formats and printed.

The \LaTeX source that is used in this book is taken from *Think Python* book.

<http://www.greenteapress.com/thinkpython/thinkpython.html>

Preface

About this book

We believe in Open-Source. And to be honest we had no great goal neither idea, when we have start it to write this book. We were just amazed by the technology, so we decided at work while we're both grabbing a coffee and talking most likely about something non-sense on the couch. However, we aim to make it simple, to be good enough to help newcomers immediately, and make it open-source because we believe in open-source.

Structure of the book

Our approach of writing this book is to make it stable and good enough for any one who wants to start Ansible. From there on we want to improve and add new chapters in future, which also means anyone who want to write a chapter, or make a correction is welcome to do so. Therefore initial version will only contain several chapters, and we will try our best to add new chapters and make corrections.

That is being said: we do not believe reading hundred of pages and than magically start doing something. We think you need to get your hands dirty, and despite this might sometimes be the hard way to learn, we think it is the best way to learn.

Many books contains just to much unnecessary and detailed information. Beginners do not need the history of a technology, it is waste of time, at least this is our view. We believe if you are interested in minor details, history and philosophy behind it you can find it your own. We want it to give the fundamentals of the Ansible, and the quickest way to start working with it. That is why we targeted to make short chapters that user can start and finish, and exercises that describe the concept of Ansible step by step.

Some of the chapters are completely based on the case studies, it means,

- A task or a problem is identified
- Way to solve it is identified
- An example of how to solve the problem step by step explained
- Concept and terminology behind the example is explained

Other chapters will combine minimal amount of text and code but will not include any unnecessary information.

Contents

| | |
|--|-----------|
| Preface | v |
| 1 Automaton | 1 |
| 1.1 Continuous Integration & Delivery & Deployment | 1 |
| 1.2 Configuration Management | 2 |
| 1.3 Orchestration | 3 |
| 1.4 Ansible | 3 |
| 1.5 To Get Started | 4 |
| 1.6 Up Shot | 5 |
| 2 The Good Old Command Line Interface | 7 |
| 2.1 To Start | 7 |
| 2.2 First Ad-hoc Commands | 8 |
| 2.3 Examples | 9 |
| 2.4 To Learn More | 11 |
| 2.5 Up Shot | 11 |
| 3 Case Study : Common Post Installation Setup | 13 |
| 3.1 The Problem | 13 |
| 3.2 Requirements | 14 |
| 3.3 Playbook | 14 |
| 3.4 Troubleshooting | 17 |
| 3.5 Up Shot | 17 |

| | | |
|----------|--|-----------|
| 4 | Case Study : Deploying Web Application with Nginx | 19 |
| 4.1 | The Problem | 19 |
| 4.2 | Variables | 20 |
| 4.3 | Templates | 20 |
| 4.4 | Loops | 22 |
| 4.5 | Handlers | 24 |
| 4.6 | Putting it Together | 24 |
| 4.7 | Up Shot | 26 |
| 5 | The Inventory | 27 |
| 5.1 | Modeling Your Infrastructure | 27 |
| 5.2 | Host Selection | 28 |
| 5.3 | Modeling Shortcut | 29 |
| 5.4 | Inventory Parameters | 29 |
| 5.5 | Host and Group Variables | 30 |
| 5.6 | Up Shot | 31 |
| 6 | The Playbook Organization | 33 |
| 6.1 | Code Reuse | 33 |
| 6.2 | Roles | 34 |
| 6.3 | Up Shot | 36 |
| 7 | Case Study : Multiple Playbooks | 37 |
| 7.1 | The Problem | 37 |
| 7.2 | The Directory Structure | 37 |
| 7.3 | Dividing Common Post Installation Playbook | 38 |
| 7.4 | Dividing Web Application Playbook | 39 |
| 7.5 | Main Playbook | 41 |
| 7.6 | Inventory File | 41 |
| 7.7 | Executing Playbook | 42 |
| 7.8 | Up Shot | 42 |

| Contents | ix |
|---|-----------|
| A How to Install Ansible | 43 |
| A.1 Install From Source Code | 43 |
| A.2 Install via Package Manager | 43 |
| B Installing and Using Vagrant | 45 |
| B.1 Requirements | 45 |
| B.2 Starting Vagrant Machine | 45 |
| B.3 Increasing Number of Machines | 46 |
| B.4 Vagrant and Ansible Code | 46 |

Chapter 1

Automaton

Chaos and complexity: it is the part of our lives, this is especially true among us. By *us*, I mean developers, system administrators, software engineers or what ever title we choose or given. We did have accepted chaos and complexity soon as we have started to get interested in the field of computer science and we love it, as much as we hate it. We build: useful, sometimes fantastic software that we are proud of it. It is beautiful and amazing thing, what one can do with a computer. We have this beautiful abstraction layer and we use other existing softwares, servers, libraries, and as result we can deliver a software product that helps other people, make their life easier, more enjoyable or make some one eventually rich.

However we need to fix bugs, implement new features and eventually update it. Based on our economical and business environment the software that is being build must be delivered frequently and process of delivery must be reliable. This is the part where the **complexity** arises; suddenly you notice that every single piece of software that your application requires to function (e.g. nginx, ntp, postfix, ruby, git). Well, we could count hundred more items, but you got the point.

Most software companies, under time and financial pressure, will ignore the any sort of configuration that will ease up their life. We have done it many times; we will do it in future as well. But soon as software start to grow and you see your system-admin sweating and trying to install one virtual machine after another, one update after another: well this is the **chaos** and it is a sign that indicates time to change things and start automation.

In this book we will address the issue of chaos and complexity by introducing a powerful automation tool. To be more precise, we will take a process that is running successfully without any problem and automated. In this sense main topics we will be dealing with it will be **Configuration Management** and **Continuous Delivery**.

1.1 Continuous Integration & Delivery & Deployment

Continuous Integration & Delivery & Deployment are by nature system-design practises. These design practises are most often misused and also little bit confusing because of the jargon, but let us clear the air.

Continuous Integration : is the practice of frequently building a software, by checking out your code from code repository (e.g. Git, SVN, CVS). There are several systems which employ this practise such as Jenkins, Travis, Bamboo, CruiseControl. The idea behind this practice is to test your code often as possible (every time developer adds a new piece of code), to reduce code errors and increase productivity with very short development cycles. In this practise developer is informed immediately when tests fails, so developer can fix the problem quickly and commit/push changes to code repository.

Continuous Delivery : is the practice of ensuring that a software can be deployed to the production at any time. You can think this as an extended version of Continuous Integration. Software product in this stage can be under further tests which automated tests can not catch, such as design issues. However, the main goal is to ensure that software can be deployed rapidly. While Continuous Integration makes sure software is buildable, Continuous Integration takes the job of delivering software product, meaning deploying actually to servers.

Continuous Deployment : is the automation of the Continuous Delivery. Meaning, code that is passed through the chain of events and passed all tests deployed to the production.

Now, why these practises do matter? Because automation does decrease and eventually removes the task of configuration. To see why this is such a big deal, please answer following questions:

- How long does it take for you to deploy a software product such as web application that your company developed?
- In case some one else in your team need to make the deployment, how long they will need?
- In case your server crashes and you need to re-configure everything, how long will this take?
- How long do you spend to prepare documentation of system configuration?

Imagine that you could execute a single command and that will take care of installing dozens of dependencies and several hours of server configuration. This is all possible, it is not new, neither a rocket science, and we are hoping to show you all with Ansible to prove a point how easy it is. That is why this practises do matter, and you should very much prefer to use your time for something valuable (such as drinking coffee on couch) than making server configuration.

1.2 Configuration Management

Configuration management is a process of configuring basic OS services and enforcing their state across IT infrastructure, however this is not limited to OS services and can be extended to services such as setting up an HTTP server (e.g. Apache). The goal here is to have several blueprints based on your infrastructure, for instance having one blueprint for your web servers and one for your database servers. Let us assume that you have twenty

web servers and four database servers. The goal here is to when ever you make a change in a blueprint, such as web servers blueprint, this is applied across all web servers. So to say, all of your servers are in same state, without someone manually making change in every single of them.

1.3 Orchestration

We are way beyond the time, where we once knew where the heck is our server. Now days we have cloud, we have servers all around the globe, and we have hundreds of them. This is the point where orchestration comes into play. Orchestration describes the process where group of automated tasks are executed, coordinated and managed. Why the fancy word? To begin with, it sounds good, on the other hand, this fancy word and concept will reduce all the errors that you are making while updating your software stack and watching TV at the same time. it will give you automation capability way beyond imagination.

1.4 Ansible

1.4.1 What is Ansible?

Ansible is open-source automation and orchestration engine. In other words, it includes all necessary IT automation tools. It is a very simple yet powerful tool to use: it does not requires hours long setup for managing your servers, it is scalable, it is secure, it is easy to understand.

1.4.2 How does it work?

Ansible works by connecting your remote servers via SSH, and will push out small programs which are executed by Ansible. These small programs will be the desired state of the infrastructure configuration. Ansible then executes these programs over SSH and deletes them when finished.

1.4.3 Why SSH?

Using SSH brings many features out-of-the-box and solves bucket-load of problems. Here are few of them:

- Agentless : No server-side software or agent required, you only need OpenSSH daemon.
- Secure: It is a very secure protocol.
- Resource Utilization: When Ansible is not managing remote server there are no resources consumed on those servers.
- Turn-Key: Allows you to start managing your machines immediately.
- Credential Segregation: It uses existing SSH credentials, as well as your privileges. For instance, if a developer is not allow to access to production server he will not be able to push contents to that server.

1.4.4 How to use it?

Using Ansible is pretty straight forwards actually, however you should posses some knowledge of :

- SSH and basic bash
- Working with the command line
- Perform basic maintenance tasks

We think at least you should have some of the knowledge that is in the scope of The Linux Professional Institute Certification(LPIC-1) first level. Read about LPIC at

http://en.wikipedia.org/wiki/Linux_Professional_Institute_Certification

1.5 To Get Started

We did not want it to have installation details of tools to became main body of the book, therefore we have added these kind of information into appendices. Under each section, you will find relevant information that you will need it, which in case of installation details you will be referred it to relevant appendix.

1.5.1 How to install it?

Please refer to *Appendix A How to Install Ansible* for installation details. If you have already installed Ansible you can skip this section.

1.5.2 Using a Vagrant Machine

This section is for those who does not have any idea about how to use virtual machines.

You need to get your hands dirty, so you need an environment to test things you learn. Therefore we encourage you to start testing while you are reading it. To solve to test-server problem, we suggest you to use Vagrant. Vagrant is a tool for creating, configuring, managing virtual machines¹. We have prepared a simple Vagrant file which you can use for firing one or several virtual machines quickly.

Please refer to *Appendix B Installing and Using Vagrant* for installation details. If you have already installed Vagrant and you know how it works you can skip this section.

1.5.3 Where are the source files?

You can find the source code of examples that are given in this book in the same repository as this book:

<https://github.com/AutomationWithAnsible/book>

Under folder *source-code* you can find the examples that corresponds to chapter number and name.

¹<http://www.vagrantup.com/>

1.6 Up Shot

In this chapter we have introduced software delivery and configuration practises and we have introduced Ansible:

- Continuous Integration : is the practice of frequently building a software and running tests against your source code to validate the functionality that has been developed.
- Continuous Delivery : is the practice of deploying software rapidly and automatise fashion without any developer involvement.
- Continuous Deployment : is the automation of the Continuous Delivery.
- Configuration Management : is a process of configuration of basic OS services and enforcing their state across IT infrastructure.
- Orchestration : is a process where group of automated tasks are executed, coordinated and managed.
- Ansible : is an open-source IT automation and orchestration engine.

Chapter 2

The Good Old Command Line Interface

We have this senseless discussion, whether Graphical User Interface(GUI) or Command Line Interface(CLI) is better. At the end of the argument we both agree that they have different purposes and somethings just does not make sense doing with GUI and other things with CLI. But we get back to discussion quite often. Long story short, CLI is handy, especially for a System Administrator or a Developer.

In this chapter we will introduce Ansible Ad-Hoc Commands, which is a tool that might be useful for running some quick arbitrary commands. The actual goal is to get you up to speed and start using Ansible as well as explaining some basics.

2.1 To Start

Along this book, when ever you execute a code and you have a connection problem, first thing to check is your SSH connection. Ansible relays on SSH so it is better to understand it and figure out why it is not connecting. Most of connection issues are:

- Network issues : connectivity or DNS problem
- SSH connection issues (e.g. port, username/password, key)

A good way to find what is the error is to eliminate them one by one, therefore try to:

- Ping your target hostname or IP to see whether you get any respond
- If your target is responsive use normal ssh to verify your access
- Check SSH username that you are connecting. Common problem is that you are accessing the system with wrong SSH user, for instance some system prohibits root ssh user login.
- Check whether your server requires key or password based authentication

- Check whether your user have sudo rights, it might be required based on type of operation
- Last, you can pass the `-vvvv` argument to ansible which will print out some useful info including connection issues

If you have regular SSH connection from terminal that it is great, you are ready to go.

2.2 First Ad-hoc Commands

Well, it is a command line argument, so it is not that complicated. Depending on your installation ansible will have executable file called *ansible*, which is the program you can use for executing ad-hoc commands. For instance, here is a really basic example of how to ping a server.

```
ansible -i 192.168.56.150, -m ping all
```



Note

IP address in this example is a server that you need to connect. If you do not have any server that you can test please refer to *Appendix B Installing and Using Vagrant*

In the example above `-i` means inventory, this is where you specify your target server. You can also specify more than one server by adding them after comma (.). The `-m` stands for module name, meaning module that will be executed is the *ping* module. We will be using modules a lot and in next chapter we will already start to dive in with more detail. The *all* refers to all hosts, meaning command will be execute on all target servers, but we are using only one host for now. When the code above is executed you will see the following response from your command line.

```
192.168.56.150 | success >> {  
    "changed": false,  
    "ping": "pong"  
}
```

Figure 2.1: Response From First Ad-hoc Command

Assumption is that you are using standard SSH port (22), if you have anything other than that you can specify that as well after IP address:

```
ansible -i 192.168.56.150:24350, -m ping all
```

When you are executing ad-hoc commands by default *command* module is selected, which you do not need to declare, this module executes commands on remote servers. You can see that from following example. When you execute an ad-hoc command you can pass arguments to the module. For this purpose you can use the *-a* and give the arguments which will pass to module afterwards. In this case we want *whoami* command to be executed on server so we give that as an argument.

```
ansible -i 192.168.56.150, -a "whoami" all
```

And response we get from server will be:

```
192.168.56.150 | success | rc=0 >>
enginyoyen
```

Figure 2.2: Server Response to *whoami* Command

There are several other arguments you can add while executing *ansible* command, instead of giving example of each of them, we will explain them briefly but only relevant ones. Short and long names are giving together:

-m, --module-name: Module name to execute

-a, --args: Module arguments

-s, --sudo: Run operations with sudo

-k, --ask-pass: Ask for SSH password

-K, --ask-sudo-pass: Ask for sudo password

-u, --user: Connect remote target with different user name

-f, --forks: Specify number of parallel processes to use when communicating with remote hosts (default is 5)

2.3 Examples

Following examples are relatively minor and do not bring anything newer than what we have explained above, however it is good practice to start with them to get use to it. If you already feel comfortable with this, you can skip it to next section already.

2.3.1 Copying File

For copying file you will use *copy* module, this module will copy a file from source (your machine) to a destination (provisioned machine).

```
ansible -i 192.168.56.150, -m copy -a  
'src=myfile dest=/tmp/mynewfile.txt' all
```

In this example, as you can see argument that is passed to the module are parameterized. Meaning, source and destination of the file are being passed to the module. What this tell us is that, we can send more than one parameter to the module. For example, it is possible to set the file owner and mode.

```
ansible -i 192.168.56.150, -m copy -a  
'src=myfile dest=/tmp/mynewfile.txt mode=644 owner=jane' all
```

2.3.2 Managing Packages

Following example, will demonstrate usage of *apt* module, which will make sure that latest version of nginx HTTP server is installed.

```
ansible -i 192.168.56.150, -s -m apt -a 'pkg=nginx state=latest' all
```

Following example, will update the cache meaning it will resynchronize the package index.

```
ansible -i 192.168.56.150, -s -m apt -a 'update_cache=yes' all
```

2.3.3 Downloading Files

Following example, will demonstrate usage of *get_url* module, which downloads a file from remote server to your target server. It support HTTP, HTTPS, and FTP protocols. In this example we will give two parameters, url, which is the URL path of the file that we want to download, and dest which is the location where should be saved in target host.

```
ansible -i 192.168.56.150, -s -m get_url -a 'url=http://apache.imsam.info/  
tomcat/tomcat-8/v8.0.3/bin/apache-tomcat-8.0.3.tar.gz  
dest=/tmp/apache-tomcat-8.0.3.tar.gz' all
```

Additionally you can also give SHA-256 checksum of the file, to make sure file is downloaded correctly. Which is a useful feature for many file operations.

2.4 To Learn More

It is great to use the modules, and you will be using them for everything when you are writing playbooks. However, you can not learn all the modules, and all the parameters they accept, but you can find it out easily. To get detailed account of an module you can use *ansible-doc* command and name of the module. When you execute the following command you will get complete the documentation of the module.

```
ansible-doc copy
```

In Figure 2.3 on page 11 you can see the partial output of this command.

```
> COPY

The [copy] module copies a file on the local box to remote
locations.

Options (= is mandatory):

- backup
    Create a backup file including the timestamp information so
    you can get the original file back if you somehow clobbered it
    incorrectly. (Choices: yes, no)

- content
    When used instead of 'src', sets the contents of a file
    directly to the specified value.

= dest
    Remote absolute path where the file should be copied to. If
    src is a directory, this must be a directory too.

- force
    the default is 'yes', which will replace the remote file when
    contents are different than the source. If 'no', the file
    will only be transferred if the destination does not exist.
    (Choices: yes, no)
```

Figure 2.3: Ansible Documentation Output

One other command that may help you to speed up is to learn available modules:

```
ansible-doc -l
```

2.5 Up Shot

In this chapter, we have introduced the ad-hoc commands which laid down the fundamentals for working with playbooks. Now you should have understanding of executing different modules, how to find them and use their parameters. This information will be quite valuable in next chapters.

Chapter 3

Case Study : Common Post Installation Setup

To begin with, we are going to do very simple exercise, which we will take over some of the basics from last chapter and develop on top of it. Furthermore, we will explain some of the key concept of Ansible in this chapter.

3.1 The Problem

Here is our problem:

We have a fresh installed *Ubuntu Precise 12.04*. Now, the first thing we need to do after installation is to upgrade and patch the system. This is a good practice because the system will be up-to-date, meaning fewer bugs, less security issues. Afterwards we want to change the hostname, because it will be easy to identify to server. Next step is to install Network Time Protocol (NTP); this will help our system time to be synchronized.

In other words, here are the steps:

- Update the apt repository: Updates the package management repository
- Upgrade the packages: Upgrades the existing software packages (e.g. security, apache), removes some of the packages and keep dependency
- Upgrade the kernel
- Change the host name
- Install NTP

This is a relatively simple example of **Common Post Installation Setup**, which can be extended to any range, but we will keep it simple for now.

3.2 Requirements

For completing this chapter following list should be completed;

- Server that is up and running
- Your public SSH key installed on remote server
- Ansible is installed on your computer



Note

If you do not have any server that you can test, please refer to *Appendix B Installing and Using Vagrant*

3.3 Playbook

In problem description, we have showed a road map of how to accomplish the tasks, one by one. Under normal circumstances what you will do is, you will complete the tasks with certain order or you will put them all in a some bash script file and execute it. It will definitely do the trick, but it won't be enough. This is where Ansible comes into play; Ansible introduced a concept called **playbooks**, which is a simple text file that describes configuration, a policy or a set of steps which can be used to configure a remote server. Playbooks are very powerful and very simple; they are expressed in YAML format, which is human-readable data serialization format ¹.



Note

A fair warning: YAML relies on outline indentation for structure which makes sometimes hard to type things quick and dirty way. Therefore we suggest that you use a good text editor with auto-indentation support, otherwise you will be spending a lot of time correcting indentation manually.

Now, let's create our first playbook. Create a new text file, and save it under the name *common-post-installation.yml*, under a directory of your choice. For Ansible, most of the playbooks starts with a list, and each item in the list is a list of key/value pairs. Here is our first playbook, that contains just a one play:

¹<http://en.wikipedia.org/wiki/YAML>


```
---
# This playbook play all post common installation for all servers

- hosts: all
  sudo: True

  tasks:
    - name : Update apt cache
      apt :
        update_cache=yes

    - name : apt upgrade
      apt :
        upgrade=yes

    - name : apt dist
      apt :
        upgrade=dist

    - name: Change hostname
      hostname:
        name=web01

    - name: install ntp
      apt:
        pkg=ntp
        state=present
```

The three dash at the top of the file indicates the beginning of the file, this is YAML syntax. The **hosts** line holds the target group, meaning which servers we want this playbook to run. In each playbook, there can be more than one play, meaning group of hosts can be tied up to certain tasks. For the sake of simplicity, we gave the value **all** to parameter **hosts**, in next chapter we will use other examples, where we will be using hosts declaration.

The **tasks** line holds the list of tasks that we want it to be accomplish in order and they are executed in order, one at a time, against all machines that are specified. Each task has a name and job to execute a module with a specific argument. The name in each task will be used in the output, it is a good practice to give names, we suggest that you do it as well. In our playbook we are using two modules *apt* and *hostname*, which manages apt-packages and changes the hostname respectively. In this example we use different arguments against same module, such as updating cache and upgrading. This arguments are very well documented, which in the following chapters we will get back to it again.

Now it is time to execute the tasks. For executing playbooks we will be using *ansible-playbook*.

```
ansible-playbook -i 192.168.56.150, /[FilePath]/common-post-installation.yml
```

Soon as you execute the code, you should see that ansible start working, and you will see the output. Output will give you details of the tasks whether succeeded, any-change has been made or failed. If one of the tasks is failed while running a play, you can solve the problem and re-run the play. Here is how the console output should look like:

```
PLAY [all] *****

GATHERING FACTS *****
ok: [192.168.56.150]

TASK: [Update apt cache] *****
ok: [192.168.56.150]

TASK: [apt upgrade] *****
changed: [192.168.56.150]

TASK: [apt dist] *****
changed: [192.168.56.150]

TASK: [Change hostname] *****
changed: [192.168.56.150]

TASK: [install ntp] *****
changed: [192.168.56.150]

PLAY RECAP *****
192.168.56.150      : ok=6    changed=4    unreachable=0    failed=0
```

Figure 3.1: Common Post Installation Output

In case of re-running a play, all tasks will run from top to bottom. As we have explained before, the goal is to bring a system to the desired state, in re-run, tasks that were successful **will-not** make any changes in the system again, meaning they will only make the changes that are required to bring the system to the desired state. This allows multiple execution of same play. To try out this, what you can simply do is, after your play runs once successfully, re-execute the command and watch the output. Result you get will be OK, but there won't be any changes. This is called **idempotent**, meaning change commands are not run unless needed ². So to say, every-time you execute this play, system will not install NTP.



Note

You could write all your modules arguments in one line as you can see in this example. However code format that we will follow in this book is to use new line for every argument as in previous play.

```
- name: install ntp
  apt: pkg=ntp state=present
```

²<http://en.wikipedia.org/wiki/Idempotence>

3.4 Troubleshooting

- No response from Ansible or stuck after *GATHERING FACTS* message.
Most likely hanged while waiting for sudo password, to solve it, add following argument

`--ask-sudo-pass or -K`

and it will ask the sudo password.
- SSH encountered an unknown error during the connection message
Check your server whether is running or not, check your connection via terminal to establish a SSH connection

3.5 Up Shot

In this chapter, we have made a very simple example, which introduce the following concepts:

- Tasks : List of actions that you want to be taken in remote server
- Playbook : Configuration management file, that holds all relevant information to achieve the desired state on your system
- Modules : A tool to accomplish a certain task

For finding source code of this example please refer to *subsection 1.5.3 Where are the source files?*

In next chapters we will dig into all these concepts one by one with more details and examples.

Chapter 4

Case Study : Deploying Web Application with Nginx

In this chapter we will show very basic example of how you could deploy an simple web application. Furthermore, we will explain some basic concepts such as variables, loops, templates and handlers that will be very useful for future playboo deployment.

4.1 The Problem

Here is our problem:

We do have a PHP web application and we want to deploy on a server. Naturally, we need a web server which will be the nginx, this also means there will be some server configuration.

In other words, here are the steps:

- Install nginx web server
- Install PHP-FPM
- Apply configuration
- Copy web application to remote machine
- Restart nginx

So it is a relatively simple task, but do not worry, we will introduce some new things as well.

4.2 Variables

It is most likely that you will face a situation that some of your server should behave or do something different than others despite they might be all classified as web servers. For instance, you may have a bunch of development servers which one of them is public and you want some special rights to be applied for that server. In such a case setting variables for particular host or a group might be really useful. These custom variables can be used easily by playbooks. In our example, we will have two configuration files, one of them is nginx configuration file, and the other one default configuration file of the web application. Both of these files will have some variables, that we want it to replace it. For time being, we will define variables in playbooks. There are other ways to do that, but for now we will take a simple approach. In playbook, you can define the variable like this:

```
---
- hosts: all
  sudo: True
  vars:
    delete_default_vhost: false
    user: www-data
    worker_processes: 4
    pid: /var/run/nginx.pid
    worker_connections: 768
    httpPort: 80
```

As you see it is relatively simple, everything under *vars* are treated as variable and available to be used for templates and tasks.

4.3 Templates

As we have mentioned, we have two configuration files, which we want these variables to be injected in them. For this purpose Ansible uses Jinja2 templating system, which is a designer friendly templating language for Python, modelled after Django's templates¹. For referencing a variable you have to put double curly braces around the variable name. Here is a simple example of how should it look like:

```
{{worker_connections}}
```

4.3.1 Template Files

We will be using variable references in *.j2* file(a text file) which we will deliver to the server. In this case we will have a template file, which will be structured in a way that we want it to be deliver. Values that we do want to be dynamic will be replaced with variables names with double curly braces around. As you can see from following example, variables that are defined in playbook are used in template file. Here is the nginx configuration file, called *nginx.conf.j2*:

¹<http://jinja.pocoo.org/docs/>

```

user {{ user }};
worker_processes {{ worker_processes }};
pid {{ pid }};
events {
    worker_connections {{ worker_connections }};
}

http {
    # Basic Settings
    sendfile on;
    tcp_nopush on;
    tcp_nodelay on;
    keepalive_timeout 65;
    types_hash_max_size 2048;

    include /etc/nginx/mime.types;
    default_type application/octet-stream;

    # Logging Settings
    access_log /var/log/nginx/access.log;
    error_log /var/log/nginx/error.log;

    # Virtual Host Configs
    include /etc/nginx/conf.d/*.conf;
    include /etc/nginx/sites-enabled/*;
}

```

So what we accomplish here is having a static file, which content will be determined based on variables that are injected in. And here is the web application configuration file, called *default.j2*:

```

server {
    listen    {{httpPort}};
    listen    [::]:{{httpPort}} default ipv6only=on;

    root /usr/share/nginx/www;
    index index.html index.htm;

    server_name localhost;

    location / {
        try_files $uri $uri/ /index.html;
    }

    location ~ /\.php$ {
        try_files $uri =404;
        fastcgi_pass 127.0.0.1:9000;
        fastcgi_index index.php;
        include fastcgi_params;
    }
}

```

```
}
}
```

**Note**

You do not have to understand the nginx configuration, it is just an example that we are using here.

4.3.2 Template Module

We have now templates files and variables, only thing that is left is to use the template module to transfer configuration files from source to destination. As you can see from the following example, for this purpose we are using *template* module and we specify the source file, which is the template file and the destination.

```
- name: write nginx.conf
  template:
    src=nginx.conf.j2
    dest=/etc/nginx/nginx.conf
```

Variables that are defined in template file will be replaced by real values without further work. As a result, destination server will have file called *nginx.conf* and all variables are replaced with real values.

4.4 Loops

In previous chapter we have explain how you can install a software package on a server. As you will remember it, it was relatively simple, in this example we have two dependencies, one of them is *nginx*, which is the http server, the other one is *php5-fpm*, which is the dependency required to run PHP application. If we want to write the code that will make the installation of these dependencies code should look like this (based on what we have shown in previous chapter) :

```
- name: Install nginx
  apt:
    pkg=nginx
    state=present"

- name: Install php-fpm
  apt:
    pkg=php5-fpm
    state=present"
```


This code is completely functional without any error. In this example there are only two dependencies so writing code for it is quick. However when you have dozens of them in one play duplicating line after line is not productive. Doing many similar tasks one after the other, such as package installation, user creation or copying file is a repeated tasks, and it is not efficient to copy and paste.

For this purpose you can use loops. A loop is a sequence of statements which is specified once but which may be carried out several times in succession². . Following is a standard loop, which you can specify a list with *with_items* directive and give the variable reference *item* to the *pkg* argument as parameter. This code is equivalent as one above.

```
- name: Install nginx and php-fpm
  apt:
    pkg="{{item}}"
    state=present
  with_items:
    - nginx
    - php5-fpm
```



Note

Specifying list of items is also possible with traditional array structure :

```
- name: Install nginx and php-fpm
  apt:
    pkg="{{item}}"
    state=present
  with_items:
    ["nginx", "php5-fpm"]
```

Because our web application will have some files and we have to copy them to the server, for this purpose we can use loops for copying files as well:

```
- name: Copy files
  copy:
    src="{{item}}"
    dest="/usr/share/nginx/www/{{item}}"
  with_items:
    - index.html
    - myapp.php
```

²[http://en.wikipedia.org/wiki/Loop_\(computing\)#Loops](http://en.wikipedia.org/wiki/Loop_(computing)#Loops)

4.5 Handlers

Now we know how to install nginx, copy files and make the configuration files based on variables and templates. But we also need to restart the server because we do change the configuration of the server. Despite there are many way to do this kind of actions, best way to go is to use something easy and simple: *handlers*. In ansible handlers are just list of tasks, like others, difference is that for a task to be executed handler have to be notified. Let's give an example for this: following handler will restart the nginx when it is being notified by a task. In the example below when template module completes it's job it notifies the handler. The *notify* directive contains the name of the handler.

```
tasks:
  - name: write nginx.conf
    template:
      src=nginx.conf.j2
      dest=/etc/nginx/nginx.conf
    notify:
      - restart nginx

handlers:
  - name: restart nginx
    service:
      name=nginx
      state=restarted
      enabled=yes
```

The *notify* directive registers an change event and informs the handlers. If there are no changes there are no triggers.

In the example above, *service* module, controls daemon on remote hosts. The parameter *state* in this case is set to *restart*, meaning it will always restart the nginx. The parameter *enabled* means service is enabled on boot.

The beauty of handlers is that they could be notify *n* time, but they will be executed only ones after all tasks are completed in a play.

4.6 Putting it Together

Now here is the complete play with variables, template module, loops and handlers.

```
---
- hosts: all
  sudo: True
  vars:
    delete_default_vhost: false
    user: www-data
    worker_processes: 4
```

```
pid: /var/run/nginx.pid
worker_connections: 768
httpPort: 80

tasks:
- name: update cache
  apt:
    update_cache=yes

- name: Install nginx and php-fpm
  apt:
    pkg="{{item}}"
    state=present
  with_items:
    - nginx
    - php5-fpm

- name: write nginx.conf
  template:
    src=nginx.conf.j2
    dest=/etc/nginx/nginx.conf
  notify:
    - restart nginx

- name: write my default site conf
  template:
    src=default.j2
    dest=/etc/nginx/sites-enabled/default
  notify:
    - restart nginx
    - restart php5-fpm

- name: Copy files
  copy:
    src="{{item}}"
    dest="/usr/share/nginx/www/{{item}}"
  with_items:
    - index.html
    - myapp.php

handlers:
- name: restart nginx
  service:
    name=nginx
    state=restarted
    enabled=yes

- name: restart php5-fpm
  service:
    name=php5-fpm
```

```
state=restarted
enabled=yes
```

You can execute the play with the following command:

```
ansible-playbook -i 192.168.56.150, deploying-web-application.yml
```

After you execute the command, you can go to browser and type the IP address, and check whether everything is working or not.

```
http://192.168.56.150/
```

4.7 Up Shot

In this chapter, we have made a very simple example, which introduce the following concepts:

- Variables : Values that can be replaced in tasks and can be different in each task.
- Templates : A file with placeholder values. Placeholder values are replaced by real values when it is executed by Ansible.
- Loops : A method to handle same type of task with different input and output many times.
- Handlers : List of tasks that are only executed when it is notified by an event or a change.

Chapter 5

The Inventory

We have previously mentioned that Ansible runs against your infrastructure, meaning you execute Ansible against several or hundred servers but you do the execution at once. If you think about it, the goal is to lower your workload and automatize everything, for this you need to know how many server you have, what are the address, and what type of server they are, meaning what a particular server do. In Ansible, for this purpose host files, or inventory files are being used. The used file format is INI, which we are sure many of you have heard about it, is an informal standard for configuration files for some platforms or software. INI files are simple text files with a basic structure composed of sections and properties¹.

5.1 Modeling Your Infrastructure

That being said, here is an example of an inventory file, which consist of three section called **webservers**, **dbservers** and **devservers**. Each group has a properties which consist of host alias and host address. Please examine the following inventory file shortly before you go forward. From the example you can see that, you can use one host in more than one group. In this example, there are two server that is used for development purpose but they are still part of web and database server as well.

```
[webservers]
tom001 ansible_ssh_host=tom001.myserver.com
tom002 ansible_ssh_host=tom002.myserver.com
tom003 ansible_ssh_host=tom003.myserver.com
tom004 ansible_ssh_host=tom004.myserver.com

[dbservers]
db001 ansible_ssh_host=db001.myserver.com
db002 ansible_ssh_host=db002.myserver.com
```

¹http://en.wikipedia.org/wiki/INI_file

```
[devservers]
db002 ansible_ssh_host=db002.myserver.com
tom004 ansible_ssh_host=tom004.myserver.com
```

As we have said, inventory files hold the information of your servers. This information is structured based on your infrastructure and definition of the servers, meaning you may have a set of web servers and set of database servers, and it is more likely that your inventory file will have the similar sections. The reason why your inventory file will reflect your infrastructure is because of their inner relationship. For instance : you may create a playbook which is only applicable to your database servers and one playbook only applicable to your web servers. When Ansible executes a playbook, it selects a portion of inventory file which in this case could be database servers, web servers or both.

By default inventory file is located under `/etc/ansible/hosts`, meaning hosts that are declared in this file will be available to Ansible in execution time unless otherwise is stated. However, in this book we will follow the tradition of specifying the host file when we execute a playbook. In previous examples we have always specify one IP address as host target, but with the inventory file in place we can change this with the inventory file name. For instance: if we want to execute common post installation setup against all servers we will be using following code:

```
ansible-playbook -i /[FilePath]/hosts.ini
                  /[FilePath]/common-post-installation.yml
```

**Note**

You do not need to add comma(,) when you refer a file as host argument.

5.2 Host Selection

So far we have used *all* at the end of each command, other possibility is to select a portion of host file or single alias. We will be using the host file example from section 5.1 Modeling Your Infrastructure. In this matter, if we want to use ping module against *web servers*, *db servers* and *dev servers* groups we will do it like this:

```
ansible -i hosts.ini -m ping all
```

Code above will execute ping command against all servers and will print out results. You could also select single *web servers* host:

```
ansible -i hosts.ini -m ping tom001
```

To select all hosts that matches pattern *tom**, this will select tom001 till tom004, which is four hosts:

```
ansible -i hosts.ini -m ping tom*
```

Probably most common and usable one is to select whole section, following example selects all hosts under *webservers*:

```
ansible -i hosts.ini -m ping webservers
```

5.3 Modeling Shortcut

It is great that you can just put all your hosts in one or several file and start working with it. However it is too much work, most people will have structural names with number as shown above and it does not make sense to write all of them one after the other. It is in ordered list anyway, why not use shortcut? Well the shortcut is to use numeric ranges in brackets, following example represents the same as above. Following example, will correspond to four servers.

```
[webservers]
tom[001:004]                                ansible_ssh_host=tom001:004.myserver.com
```

The ranges that is specified can also be alphabetic.

```
tom[a:d]                                    ansible_ssh_host=tom[a:d].myserver.com
```

5.4 Inventory Parameters

In the example that is given above, assumption is that your ssh credentials and connection settings are all set in your ssh configuration file, for instance:

```
Host tom4
Hostname tom004.myserver.com
Port 23480
ForwardAgent yes
ForwardX11 yes
User adham
```

However, this is not always the case, furthermore there may be cases where you would prefer to add this information to your inventory file. In the following example host, port and ssh user is given as an external parameter.

```
web01 ansible_ssh_host=192.168.1.122 ansible_ssh_port=23480 ansible_ssh_user=adham
```

If your SSH port is different following example can also be used:

```
tom[001:004].myserver.com:23480
```

For system that has more than one python, ruby or perl interpreter you can choose specific version of interpreter for that host.

```
tom[001:004].myserver.com ansible_python_interpreter=/usr/local/bin/python
```

If you are using multiple SSH keys and you do not want to use SSH agent you can specify private key file.

```
tom[001:004].myserver.com ansible_ssh_private_key_file=/Users/adham/.ssh/prod.pem
```

5.5 Host and Group Variables

In chapter 4 Case Study : Deploying Web Application with Nginx, we have introduced how to define and use variables. Beside playbooks, it is also possible to define variables in inventory files. So here is an example of simple host variable called *user* and *maxConnections* which is an arbitrary variable and maximum amount of open connections.

```
[webservers]
tom[001:004].myserver.com user="John Doe"

[devservers]
dev1.myserver.com maxConnections=1000
```

Variables that are defined in inventory files can also be used in playbooks same way of using other variables.

In the example above, all the servers under *webservers* group will have the same variable. There is a different method for this as well which is called group variable, that is applied to all group children. Following example, illustrates it:


```
[webservers]
tom[001:004].myserver.com
```

```
[webservers:vars]
person="John Doe"
```

In this case all hosts under section *webservers* will have the variable *person* assigned.

5.6 Up Shot

In this chapter we have explain what inventory is and how you can use it to ease up your work-flow. There is also possibility of generating your inventory file dynamically, which we will explain later details(This chapter is not planned for fiirst release of this book).

- Inventory File : File that holds your system informations such as server name, address, port, ssh user.
- Inventory File Type : INI file types are used for inventories, which is a simple text file that consist of sections and properties.
- Inventory Variables : A way to assign variable to a specific server or group of servers.

Chapter 6

The Playbook Organization

We are all aware how quickly things can get out of hand when we do not organize our code. Matter of organizing code is not just to make the code readable or easy to find it also means easy to manipulate and make it much more easy to reuse. This is why before we go any further we want it to explain the how you can organize your playbooks.

6.1 Code Reuse

Some of the system tasks can be reused for other systems as well. A good example for would be the one that we have shown in *chapter 3 Case Study : Common Post Installation Setup*. Example was a relatively simple play but we do not want to duplicate the code every time we need to update cache or install NTP. In such a situation, the method we follow should be including playbook that we have already written into another playbook as a task. Here is a reduced example, called *common-post-installation-task.yml* which is saved under directory *tasks*.

```
---
tasks:
  - name : Update apt cache
    apt :
      update_cache=yes

  - name : apt upgrade
    apt :
      upgrade=yes

  - name : apt dist
    apt :
      upgrade=dist

  - name: install ntp
    apt:
```

```
pkg=ntp
state=present
```

As you can see, we do not have any host information, these are just tasks. To include this play into a playbook as a task you can use *include* directive. Here is the example of how this is done.

```
tasks:
- include: tasks/common-post-installation-task.yml
```

The same idea can be used for handlers as well. In chapter *chapter 4 Case Study : Deploying Web Application with Nginx*, we have gave an example of handlers, and in here we want to show how we can reuse same handler again. In this case we will strip the handlers into one file, and save it as *common-handlers.yml* under *handlers* directory.

```
- name: restart nginx
  service:
    name=nginx
    state=restarted
    enabled=yes
```

And here is how you can include the handlers.

```
handlers:
- include: handlers/common-handlers.yml
```

You can also include other playbooks into top level playbooks. For instance:

```
---
- hosts: all
  sudo: True

- include: server-configuration.yml
- include: commom-taks.yml
- include: webserver.yml
```

6.2 Roles

Roles are logical directory structure which ansible knows about it and work against it, meaning loading tasks, files, handlers and so on will be done based on known structure.

Using roles will remove all path related issues as well as will bring structure to your playbook. It will also reduce errors, and increase the usability of playbooks. Each role is meant to accomplish a certain set of tasks, but it is configured in a way which can be used immediately. So to say : *convention over configuration*¹ .

6.2.1 Directory Structure

In the following example there are two roles which are called *common* and *webservers*, these roles are added to the play like this:

```
---
- hosts: webservers
  roles:
    - common
    - webservers
```

Before we go further, please refer to figure 6.1 on page 35 to examine directory organization of roles.

Under each role, there are sub directories which corresponds to certain responsibility. Lets take the example of *common* role, so when, *main.yml* exist under following directories:

- roles/common/defaults
- roles/common/task
- roles/common/handlers
- roles/common/vars
- roles/common/meta

these files will be added to the play. Meaning you do not have to specify paths for tasks, handlers, variables, meta data. To make it little bit more easy, you do not have to use *include* directives for any of these files.

When you want to use files that are under following directory, *copy* and *script* tasks will make the reference to the following directory:

- roles/common/files

meaning any file or script under this directory will be available copy and script tasks automatically without having to path them relatively or absolutely. In same matter, files located under template directory will be available to template tasks, and files located under task directory will be available for include tasks.

¹http://en.wikipedia.org/wiki/Convention_over_configuration/




Figure 6.1: Role Directory Structure

- roles/common/tasks
- roles/common/templates

6.2.2 Creating Roles

Ansible offers simple way to create complete directory structure for roles. Following command will create role called *common* and empty YAML files in directories.

```
ansible-galaxy init common
```

6.2.3 Using Community Roles

Ansible Galaxy, is a free site for downloading community developed Ansible roles, which means for most common task you do not have to write any role and can reuse existing role. Please refer to:

```
http://galaxy.ansible.com/
```

6.3 Up Shot

In this chapter we have gave brief overlook of playbook organization.

- Roles : are logical directory structure which ansible knows about it and work against it. It follows the method of convention over configuration.

Chapter 7

Case Study : Multiple Playbooks

In this chapter we will put all the things we have tried so far in previous chapters in a good structural way so we can take advantage of roles in Ansible.

7.1 The Problem

Here is our problem:

in chapter 3 Case Study : Common Post Installation Setup and chapter 4 Case Study : Deploying Web Application with Nginx we have made two separate tasks, which was great. However, we want to merge these playbooks and configure it in much more friendly way. Furthermore we want to use the playbook against several server with some dynamic configuration.

In other words, here are the steps:

- Customize previous code
- Make it installable to several server
- Make the playbook much more dynamic

7.2 The Directory Structure

We have now two separate playbook which we want to join, one is common post installation and the other one is nginx installation. So what we want to do is make a role out of each playbook. Therefore we create the following directories first.

- roles/common
- roles/webservers

If you can remember it, common post installation playbook did not include any templates, scripts, neither variables. So it consisted only of tasks. Nginx installation on the other had all these extra files. So what we will do is create the directory structure and add main.yml file as default to required directories:

- roles/common/tasks/main.yml
- roles/webserver/files
- roles/webserver/templates
- roles/webserver/handlers/main.yml
- roles/webserver/tasks/main.yml
- roles/webserver/vars/main.yml

7.3 Dividing Common Post Installation Playbook

Common post installation, had tasks of updating apt cache, installing ntp and changing the hostname. If we want some logical structure out of it, what we can do is, to separate hostname task and others. In such a case the structure of the files and their content will look like this:

File : /roles/common/tasks/hostname.yml

```
---
- name: Change hostname
  hostname:
    name={{host_hostname}}
```

File : /roles/common/tasks/apt.yml

```
---
- name : Update apt cache
  apt :
    update_cache=yes

- name : apt upgrade
  apt :
    upgrade=yes

- name : apt dist
  apt :
    upgrade=dist

- name: install ntp packages
```



```
apt:
  pkg=ntp
  state=present
```

And finally we can include these two files into main.yml file so by default when common role is executed these tasks will be included in.

File : /roles/common/tasks/main.yml

```
---
- include: hostname.yml
- include: apt.yml
```

7.4 Dividing Web Application Playbook

In this case we have several different components which we have to place them in correct directories. To begin with, we need to move j2 templates to templates directory, and HTML and PHP files into files directory.

- roles/webservers/files/index.html
- roles/webservers/files/myapp.php
- roles/webservers/templates/default.j2
- roles/webservers/templates/nginx.conf.j2

Next thing we can separate is handlers. So we should put our handlers under roles/webservers/handlers/main.yml

```
---
- name: restart nginx
  service:
    name=nginx
    state=restarted
    enabled=yes

- name: restart php5-fpm
  service:
    name=php5-fpm
    state=restarted
    enabled=yes
```

As you know we also had some variables, so we can put all the variables under roles/webservers/defaults/main.yml

```
delete_default_vhost: false
user: www-data
worker_processes: 4
pid: /var/run/nginx.pid
worker_connections: 768
httpPort: 80
```

Only thing that is left is now to separate tasks. We will separate tasks in two section one is nginx installation and configuration and another one is deployment of the web application.

File :roles/webserver/tasks/setup_nginx.yml

```
---
- name: Install nginx and php-fpm
  apt:
    pkg="{{item}}"
    state=present
  with_items:
    - nginx
    - php5-fpm
  notify:
    - restart nginx

- name: write nginx.conf
  template:
    src=nginx.conf.j2
    dest=/etc/nginx/nginx.conf
  notify:
    - restart nginx

- name: write my default site conf
  template:
    src=default.j2
    dest=/etc/nginx/sites-enabled/default
  notify:
    - restart nginx
    - restart php5-fpm
```

File :roles/webserver/tasks/deploy_webapp.yml

```
---
- name: Copy files
  copy:
    src="{{item}}"
    dest="/usr/share/nginx/www/{{item}}"
```

```
with_items:
  - index.html
  - myapp.php
```

And finally we can include these two files into main.yml as we did before:

File :roles/webservers/tasks/main.yml

```
---
- include: setup_nginx.yml
- include: deploy_webapp.yml
```

7.5 Main Playbook

Main playbook will be relatively small, as it only includes host definition and roles. This file which we will be calling *webervers.yml* will be same directory as roles and here is the content of it:

```
- hosts: web
  sudo: true
  roles:
    - common
    - webservers
```

As you can see from playbook, it will run against all the host that are called web. And first it will execute common role and then webservers. We could additionally create one more playbook in same level which only executes common role, because there might be cases where we may need only common post installation and this can run against *all* servers:

```
- hosts: all
  sudo: true
  roles:
    - common
```

7.6 Inventory File

As one of the goal of this chapter was to run this ansible code against several machines, what we want to do is create a INI file which holds the hosts files. Here is the how host file looks like, and each host property has variable called *host_hostname*. This variable is used in under common post installation, when we change the host name.

```
[web]
web01      ansible_ssh_host=192.168.56.150      host_hostname=web01
web02      ansible_ssh_host=192.168.56.151  host_hostname=web02
web03      ansible_ssh_host=192.168.56.152  host_hostname=web03
```

**Note**

If you are using vagrant file that we have provided you can increase number of servers by increasing a variable in vagrant file. For details please refer to *Appendix B Installing and Using Vagrant*

7.7 Executing Playbook

You can execute the playbook same way as we have done before, however this time instead of giving server address, we can specify inventory file:

```
ansible-playbook -i hosts.ini webservers.yml
```

This will execute webservers role for all three servers.

7.8 Up Shot

In this chapter we have illustrated how you can organize playbooks by separating tasks, variables, handlers as well as running a playbook against several servers.

Appendix A

How to Install Ansible

Installing Ansible is relatively easy, most likely the best way is to checkout the repository and start using it. With this approach you can just pull latest changes and you can use the features immediately. But other methods are just fine as well.

A.1 Install From Source Code

To install Ansible from source:

```
$ git clone git://github.com/ansible/ansible.git
$ cd ./ansible
$ source ./hacking/env-setup
```

There are several dependency which you can install via python package manager, *pip*. If you don't have the pip first install pip:

```
$ sudo easy_install pip
```

Ansible also uses the following Python modules that need to be installed:

```
$ sudo pip install paramiko PyYAML jinja2 httpplib2
```

A.2 Install via Package Manager

Alternatively you can install via package manager that you are using:
via homebrew :

```
$ brew install ansible
```

via pip:

```
$ sudo pip install ansible
```

via apt

```
$ sudo apt install ansible
```

via yum

```
$ sudo yum install ansible
```

via pkg

```
$ sudo pkg install ansible
```

If you have any installation problems please refer to:

http://docs.ansible.com/intro_installation.html

Appendix B

Installing and Using Vagrant

Vagrant is a tool for creating, configuring, managing virtual machines. Because we want you to actually test the things we have shown in this book, we have prepared a simple Vagrant file which you can use for firing one or several virtual machines quickly. Furthermore, we have add a small playbook which will copy your public SSH key into vagrant machine so you don't have to do anything.

B.1 Requirements

To use vagrant you need to install vagrant and virtualbox, please download both and install it.

- <http://www.vagrantup.com/>
- <http://www.virtualbox.org/>

You also need a SSH key, we assume that your SSH keys under your home directory. In MacOS this should be:

```
/Users/adham/.ssh/id_rsa.pub
```

If you don't have a SSH key, please create one and then proceed.

B.2 Starting Vagrant Machine

To create a vagrant machine, open your terminal, go to *source-code/Appendix B - Installing and Using Vagrant* folder, and execute following command.

```
vagrant up
```

This command will create a new vagrant machine and will copy your public SSH key. By default first machine that is created will have the following IP address:

192.168.56.150

For more information about how to use vagrant please refer to:

<http://docs.vagrantup.com/>

B.3 Increasing Number of Machines

If you want to increase the number of machines that you are firing-up you can simply increase the *nodes* value in vagrant file. You can also change the default IP address.

B.4 Vagrant and Ansible Code

We have an ansible playbook which creates local user and copy SSH key. Here is the content of vagrant file:

```
Vagrant.configure("2") do |config|
  config.vm.box_url = "http://files.vagrantup.com/precise64.box"

  ## Nodes
  nodes = 1
  rangeofips = 149
  (1..nodes).each do |n|

    vmip = "192.168.56.#{rangeofips + n.to_i}"
    name = "vagrant0#{n}.local"
    puts "Let's talk about %s." % name
    ## Node Conf
    config.vm.define name do |cfg|
      cfg.vm.box = name
      cfg.vm.host_name = name
      ## Comment public network to disable bridge
      cfg.vm.network :public_network
      cfg.vm.network :private_network, ip: vmip
      cfg.ssh.forward_agent = true
      cfg.vm.provider "virtualbox" do |vb|
        ## headless or non headless machine
        vb.gui = false
      end

      ## Ansible Provisioning
      cfg.vm.provision :ansible do |ansible|
        ansible.playbook = "vagrant-provision.yml"
      end

      ## Debugging
```



```

        ansible.verbose = true
        ansible.verbose="vvvvv"
    end
end
end
end
end

```

And here is the ansible code:

```

---
#
# This playbook deploys your keys to the vagrant
#

- name: Provision my keys
  hosts: all
  sudo: True
  vars:
    localuser: "{{ lookup('ENV','USER') }}"
    mypassword : "$5$rounds=110000$Jm.keFgd6zfXrnrJ$ar4ns4Y/Vds32qqet19KlR3evMgRkdTjoIf3eL7z"
  tasks:
    - name: Create your local user
      user:
        name="{{localuser}}"
        home="/home/{{localuser}}"
        shell="/bin/bash"
        append="true"
        group="admin"
        comment="{{localuser}}"

    - name: check keys
      stat:
        path=~/.ssh/id_rsa.pub
        get_md5=no
        follow=yes
      register: mypubkey
      sudo_user: "{{localuser}}"
      connection: local

    - name: Putting you authorized_key
      authorized_key:
        key="{{lookup('file', '~/.ssh/id_rsa.pub')}}"
        user="{{localuser}}"
        manage_dir=yes
      when: "mypubkey.stat.exists == True"

    - name: Create password as you are not using keys
      user:
        name="{{localuser}}"
        password="{{mypassword}}"

```

```
when: "mypubkey.stat.exists == False"
```